

FleetDB

A Schema-Free Database in Clojure

Mark McGranaghan

January 8, 2010

Motivation

Walkthrough

Implementation

QA

Motivation

Motivation

optimize for agile development

Walkthrough

- ▶ Client Quickstart
- ▶ Basic Queries
- ▶ Concurrency Tools

Client Quickstart

```
(use 'fleetdb.client)  
(def client (connect))
```

```
(client ["ping"])  
=> "pong"
```

Insert

```
(client
  ["insert" "accounts"
    {"id" 1 "owner" "Eve" "credits" 100}])
```

Insert

```
(client
  ["insert" "accounts"
    {"id" 1 "owner" "Eve" "credits" 100}])
```

=> 1

Select

```
(client  
  ["select" "accounts" {"where" ["=" "id" 1]}])
```

Select

```
(client  
  ["select" "accounts" {"where" ["=" "id" 1]}])  
  
=> [{"id" 1 "owner" "Eve" "credits" 100}]
```

Select with Conditions

```
(client  
  ["select" "accounts"  
    {"where" [">=" "credits" 150]}])
```

Select with Order and Limit

```
(client  
  ["select" "accounts"  
    {"order" ["credits" "asc"] "limit" 2}])
```

Update

```
(client  
  ["update" "accounts" {"credits" 105}  
    {"where" ["=" "owner" "Eve"]}])
```

Explain (I)

```
(client  
  ["explain" ["select" "accounts"  
    {"where" ["=" "owner" "Eve"]}]]])
```

Explain (I)

```
(client  
  ["explain" ["select" "accounts"  
                 {"where" ["=" "owner" "Eve"]}]]])
```

```
=> ["filter" ["=" "owner" "Eve"]  
     ["collection-scan" "accounts"]]
```

Create Index

```
(client  
  ["create-index" "accounts" "owner"])
```


Explain (II)

```
(client
  ["explain" ["select" "accounts"
               {"where" ["=" "owner" "Eve"]}]]])

=> ["index-lookup" ["accounts" "owner" "Eve"]]
```

Multi-Write

Multi-Write

```
(client  
  ["multi-write"  
    [write-query-1 write-query-2 ...]])
```

Multi-Write

```
(client  
  ["multi-write"  
    [write-query-1 write-query-2 ...]])
```

```
=> [result-1 result-2 ...]
```

Checked-Write

Checked-Write

```
(client
  ["checked-write"
   read-query
   expected-read-result
   write-query])
```

Checked-Write

```
(client  
  ["checked-write"  
   read-query  
   expected-read-result  
   write-query])
```

```
=> [true write-result]
```

Checked-Write

```
(client  
  ["checked-write"  
   read-query  
   expected-read-result  
   write-query])
```

```
=> [true write-result]
```

```
=> [false actual-read-result]
```


Implementation

- ▶ Background
- ▶ Organization
- ▶ Key ideas

Background

- ▶ http://clojure.org/data_structures
- ▶ <http://clojure.org/sequences>
- ▶ <http://clojure.org/state>

Organization

Organization

- ▶ `fleetdb.core`: pure functions

Organization

- ▶ `fleetdb.core`: pure functions
- ▶ `fleetdb.embedded`: identity and durability

Organization

- ▶ `fleetdb.core`: pure functions
- ▶ `fleetdb.embedded`: identity and durability
- ▶ `fleetdb.server`: network interface

`fleetdb.core`

- ▶ Pure functions

fleetdb.core

- ▶ Pure functions
- ▶ Databases as Clojure data structures

fleetdb.core

- ▶ Pure functions
- ▶ Databases as Clojure data structures
- ▶ Read: $\text{db} + \text{query} \rightarrow \text{result}$

fleetdb.core

- ▶ Pure functions
- ▶ Databases as Clojure data structures
- ▶ Read: $\text{db} + \text{query} \rightarrow \text{result}$
- ▶ 'Write': $\text{db} + \text{query} \rightarrow \text{result} + \text{new db}$

`fleetdb.core` Query Planner

`fleetdb.core` Query Planner

- ▶ Implements declarative queries

fleetdb.core Query Planner

- ▶ Implements declarative queries
- ▶ Database + query \rightarrow plan

fleetdb.core Query Planner

- ▶ Implements declarative queries
- ▶ Database + query \rightarrow plan

```
["select" "accounts"  
 {"where" ["=" "owner" "Eve"]}]
```

fleetdb.core Query Planner

- ▶ Implements declarative queries
- ▶ Database + query \rightarrow plan

```
["select" "accounts"  
 {"where" ["=" "owner" "Eve"]}]
```

```
["filter" ["=" "owner" "Eve"]  
 ["record-scan" "accounts"]]
```

fleetdb.core Query Planner

- ▶ Implements declarative queries
- ▶ Database + query \rightarrow plan

```
["select" "accounts"  
 {"where" ["=" "owner" "Eve"]}]
```

```
["filter" ["=" "owner" "Eve"]  
 ["record-scan" "accounts"]]
```

```
["index-lookup" ["accounts" "owner" "Eve"]]
```


`fleetdb.core` Query Executor

fleetdb.core Query Executor

- ▶ Database + plan \rightarrow result

fleetdb.core Query Executor

- ▶ Database + plan \rightarrow result

```
["filter" ["=" "owner" "Eve"]  
  ["record-scan" "accounts"]]
```

fleetdb.core Query Executor

- Database + plan \rightarrow result

```
["filter" ["=" "owner" "Eve"]  
  ["record-scan" "accounts"]]
```

```
(filter (fn [r] (= (r "owner") "Eve"))  
  (vals (:rmap (db "accounts"))))
```

`fleetdb.embedded`

- ▶ Wraps `fleetdb.core`

`fleetdb.embedded`

- ▶ Wraps `fleetdb.core`
- ▶ Adds identity and durability

`fleetdb.embedded`

- ▶ Wraps `fleetdb.core`
- ▶ Adds identity and durability
- ▶ Databases in atoms

fleetdb.embedded

- ▶ Wraps `fleetdb.core`
- ▶ Adds identity and durability
- ▶ Databases in atoms
- ▶ Append-only log

fleetdb.embedded Read Path

`fleetdb.embedded` Read Path

- ▶ Dereference database atom

fleetdb.embedded Read Path

- ▶ Dereference database atom
- ▶ Pass to fleetdb.core

fleetdb.embedded Read Path

- ▶ Dereference database atom
- ▶ Pass to fleetdb.core
- ▶ Return result

`fleetdb.embedded` Write Path

fleetdb.embedded Write Path

- ▶ Enter fair lock

`fleetdb.embedded` Write Path

- ▶ Enter fair lock
- ▶ Dereference database atom

`fleetdb.embedded` Write Path

- ▶ Enter fair lock
- ▶ Dereference database atom
- ▶ Pass to `fleetdb.core`

fleetdb.embedded Write Path

- ▶ Enter fair lock
- ▶ Dereference database atom
- ▶ Pass to fleetdb.core
- ▶ Append query to log

fleetdb.embedded Write Path

- ▶ Enter fair lock
- ▶ Dereference database atom
- ▶ Pass to fleetdb.core
- ▶ Append query to log
- ▶ Swap in new database value

`fleetdb.embedded` Write Path

- ▶ Enter fair lock
- ▶ Dereference database atom
- ▶ Pass to `fleetdb.core`
- ▶ Append query to log
- ▶ Swap in new database value
- ▶ Return result

fleetdb.server

- ▶ Wraps `fleetdb.embedded`

fleetdb.server

- ▶ Wraps `fleetdb.embedded`
- ▶ Adds JSON client API

Aside: Source Size

Aside: Source Size

Lines of Code	
core	630
embedded	130
server	120
lint	280
utilities	140
total	1300

Takeaways

Takeaways

- ▶ Clojure's data structures are awesome

Takeaways

- ▶ Clojure's data structures are awesome
- ▶ Clojure is viable for infrastructure software

Takeaways

- ▶ Clojure's data structures are awesome
- ▶ Clojure is viable for infrastructure software
- ▶ Try FleetDB!

Thanks for listening!

Questions?

<http://FleetDB.org>

<http://github.com/mmcgrana>